

Processes Management

Sistemas de Computadores

2017/2018

Luís Miguel Pinho / Luis Ferreira / Luís Nogueira / Nuno Pereira

Processes

- ⦿ Instance of execution of a program
- ⦿ **Program != Process**
 - Program: code + static data
 - Process: dynamic instantiation of code + data + state
- ⦿ Program and process
 - Process > program: more than code and data
 - Program > process: one program can create several processes

Process

- ◉ From the point of view of the operating system, a process is the entity to which resources are given
 - CPU time, memory, files, I/O devices, etc.
- ◉ Processes are isolated from each other
 - Sequential execution has well defined interactions
 - By default, no memory sharing between processes

OS process management

- ◉ Identify and represent processes internally
- ◉ Create/remove processes of user and of system
- ◉ Schedule processes access to the CPU
- ◉ Provide communication and synchronization mechanisms

Processes identification

- ⦿ A process is identified through um unique value, names Process ID (PID)
- ⦿ PIDs are usually given by OS sequentially until the limit of the data type
 - In that case, PID numbers are recycled from older processes that have already finished (lower values first)

Process creation

- ◉ A process is created ...
 - By action of the OS
 - By action of the user
 - By another process

```
#include <sys/types.h>  
#include <unistd.h>
```

```
pid_t fork()
```

Process hierarchy

- ◉ The process that invokes the fork system call is named the *parent*
- ◉ The newly created process is the *child*
- ◉ Each one can create more processes, leading to a process tree
 - Each process has exactly one parent and zero or more children

Process hierarchy

```
$ pstree | more
```

```
init--acpid
```

```
| -avahi-daemon---avahi-daemon
```

```
| -bonobo-activati---{bonobo-activati}
```

```
| -cron
```

```
| -cupsd
```

```
| -gdm---gdm--Xorg
```

```
|      `--x-session-manag--gnome-panel---{gnome-panel}
```

```
|      |      | -gnome-settings--+-pulseaudio+-gconf-helper
```

```
|      |      |      |      `--2*[{pulseaudio}]
```

```
|      |      |      |      `--{gnome-settings-}
```

```
|      |      | -konsole---3*[bash]
```

```
|      |      | -metacity
```

```
|      |      | -ssh-agent
```

```
|      |      `--{x-session-manag}
```

```
| -getty
```

```
| -konsole--+-2*[bash]
```

```
|      | -bash---vim
```


getpid and *getppid*

```
#include <sys/types.h>
#include <unistd.h>

pid_t getpid();
pid_t getppid();
```

- ◉ *getpid* returns the PID of the process
- ◉ *getppid* returns the PID of the parent process

Process creation

- ⦿ The Child process is a copy of the parent
 - The only differences are the PID (remember it is unique), and certain resources (e.g. signals are not inherited) and stats
 - Child and parent are independent processes and both compete for the CPU time
- ⦿ After *fork*, both processes execute the same code, in the instruction immediately after the fork function call
 - Order of execution is not known
 - There needs to be a way to distinguish them

fork return values

- ⦿ -1, in case of error
 - It was not possible to create the child, so a single process still exists
- ⦿ **The PID of the child** (> 0) is returned in the parent process
 - Parent may need this PID to communicate or synchronize with the child
- ⦿ The child process receives from fork the value **zero**
 - The fork was not executed by the child
- ⦿ The different values received by parent and child allows to separate the code each one will execute

Example – fork returned values

```
pid = fork();

if(pid == -1){
    perror("Fork failed"); exit(1);
}

if(pid > 0)
    printf("Luke, I'm your father\n");
else
    printf("I am the son\n");

printf("Who executes this line?");
```

Process creation

- ◉ We can think of *fork* as a cloning function
 - It creates a copy of the original process, with the same code and data
- ◉ In the child process, the new variables have the same value as the original ones had in the parent in the instant of execution of the fork
- ◉ However, any change made in any of the processes after the fork, is independent and does not affect the other process

Example – changes after *fork*

```
int x = 10;

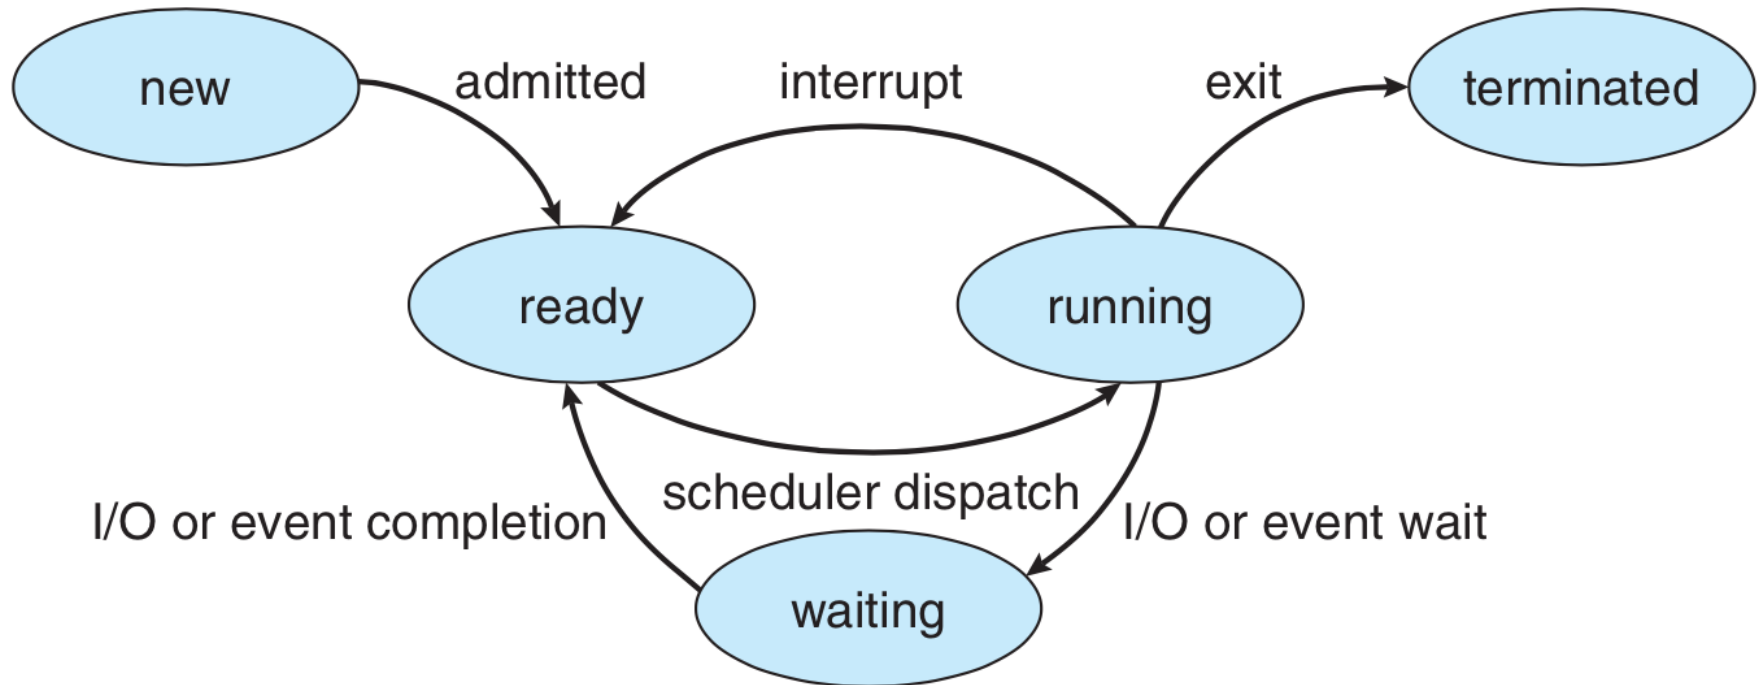
pid = fork();

if(pid > 0){
    printf("Luke, I'm your father\n");
}
else{
    printf("I am the son\n");
    x++;
}
printf("x = %d\n", x);
```

Concurrent execution

- ⦿ Only one process may, at a given instant, be executing in the core of a CPU...
- ⦿ ... but multiple processes can be loaded to memory and be executed concurrently
- ⦿ The goal of the OS is to have always a process executing, maximizing resource usage
 - Decision based in metrics of performance and load (scheduling)

States of a process



sleep

```
#include <unistd.h>
```

```
int sleep(int seconds);
```

- ⦿ Suspends execution of the process during (at least) n seconds
 - If a process receives a signal while it is sleeping, it may waken before time
- ⦿ Sleep can be used to try to obtain a order between processes
 - But it is not guaranteed, so use it only for testing, not to guarantee order and synchronization!

Process termination

- ⦿ A process terminates when ...
 - Executes the last instruction
 - Invokes the *exit* system call
 - voluntarily in an error
 - Involuntarily in an error
 - By action of a user or another process

Process termination

```
#include <stdlib.h>
```

```
void exit(int status);
```

- Terminates the process, providing a status value to the operating system
 - A single byte of *status* is used so values are 0 - 255
- Two constants are used, EXIT_SUCCESS and EXIT_FAILURE, to indicate success or failure in the execution

Process termination

- ◉ When a process terminates, the parent is informed by the OS using the SIGCHLD signal
- ◉ This event is asynchronous – the parent can receive it at any time
 - The parent can ignore the signal (the default behavior) or provide a function to be executed when the signal is received (a handler)

Zombie Processes

- ⦿ If the child terminates before the parent
 - It becomes a zombie – cannot execute but the OS still has to keep the information of its termination
- ⦿ OS cannot remove the process from the process table, but can free all process resources
 - Parent has to acknowledge the termination for the entry to be removed
 - Acknowledgement is with *wait/waitpid* functions
 - After acknowledgement, process disappears and is no longer zombie

Orphan Processes

- ◉ If the parent terminates before the child
 - Child become *orphan*
- ◉ In that case, it is necessary the process to be adopted by another process to maintain hierarchy
- ◉ In UNIX, orphan processes are adopted by the *init* process that periodicall invokes *wait* to all its children

wait and waitpid

```
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

- ◉ ***wait* blocks until any of the process children terminates**
 - Returns the PID of the terminated child, or -1 in case of error
- ◉ ***status* should be initialized with the address of a variable that will receive the exit value of the child**
 - The exit value also includes some extra information by the OS

wait and waitpid

```
pid_t waitpid(pid_t pid, int *status,  
int options);
```

- ⦿ ***waitpid* blocks until the specified child terminates**
 - *pid* = -1 means it will wait for any children, as *wait*
- ⦿ *options* may be 0 (blocks), WNOHANG (does not block) and WUNTRACED (children in the STOPPED state)

wait and *waitpid*

- ◉ When invoking *wait* or *waitpid*, the process may:
 - Be blocked until the child terminates
 - Return immediately if the child has already terminated
 - Return immediately with an error, if there are no children

Example – synchronizing with *wait*

```
pid = fork();

if(pid > 0) {
    printf("Luke, I'm your father\n");
    wait(NULL);
    printf("The child has terminated");
}
else{
    printf("I'm the son\n");
    Some_code();
    exit(0);
}
```

Process exit information

- ◉ The OS also stores some information related to the way a process terminates:
 - The exit value
 - The signal that caused the abnormal termination (if any)
 - If the process originated a *core dump*
 - etc
- ◉ There are several macros that can be used to obtain that information from the status obtained in *wait/waitpid*

Process exit information

- ◉ WIFEXITED(status)
 - Checks if the child terminated normally
- ◉ WEXITSTATUS(status)
 - Exit value (1 byte) given in the *exit* function
- ◉ WIFSIGNALED(status)
 - Checks if the child terminated due to a non-treated signal

Process exit information

- ◉ WTERMSIG(status)
 - Gets the number of the signal that terminated the child
- ◉ WCOREDUMP(status)
 - Checks if a *core dump* was generated
- ◉ WIFSTOPPED(status)
 - Checks if the child is in the STOPPED state
- ◉ WIFCONTINUED(status)
 - Checks if the process was continued

Example – process exit information

```
pid_t pid1, pid2;
int status;

pid1 = fork();
if(pid1 > 0){
    pid2 = waitpid(pid1, &status, 0);
    if(WIFEXITED(status))
        printf("Parent: child %d returned
%d\n", pid2, WEXITSTATUS(status));
    }
else
    exit(10);
}
```

Exercise

- ⦿ Implement a program that creates 2 new processes, with the following behaviour
 - Child 1 sleeps 5 seconds, prints its PID and terminates with value 1
 - Child 2 prints its PID and the parent PID and terminates with value 2
- ⦿ Independently of the order of execution and termination, parent waits first for child 1 and then child 2, printing their PIDs and exit values
 - But guarantee that both children execute concurrently