

Signals

Sistemas de Computadores
2020/2021

António Barros / Luís Ferreira

Dealing with the unexpected

- ◉ The main body of a program defines the *expected* sequence of instructions that a process should perform.
- ◉ However, unexpected events may occur during the execution of a process.
 - Such events may require special handling.
 - But if they are not predicable, WHERE should we code the handling?

Signals

- ⦿ A *signal* is a notification of an event sent to a specific process.
 - Process receives an **integer** that maps to a specific **event**.
- ⦿ Notifications are **asynchronous**.
 - A process may receive a notification at any time, during its execution.
 - If you can't predict when an event occurs, then it is not practical to handle it in the main body of the program.




- **The solution is to write specialized code to handle the signal and execute it only when necessary!**

Causes of signals

- ⦿ Hardware interrupt
 - Exceptions such as invalid memory access or a division by zero.
 - Special key combinations (e.g. CTRL-C).
- ⦿ Signals raised by the user.
- ⦿ Signals raised by other processes.
- ⦿ Signals raised by the process itself.

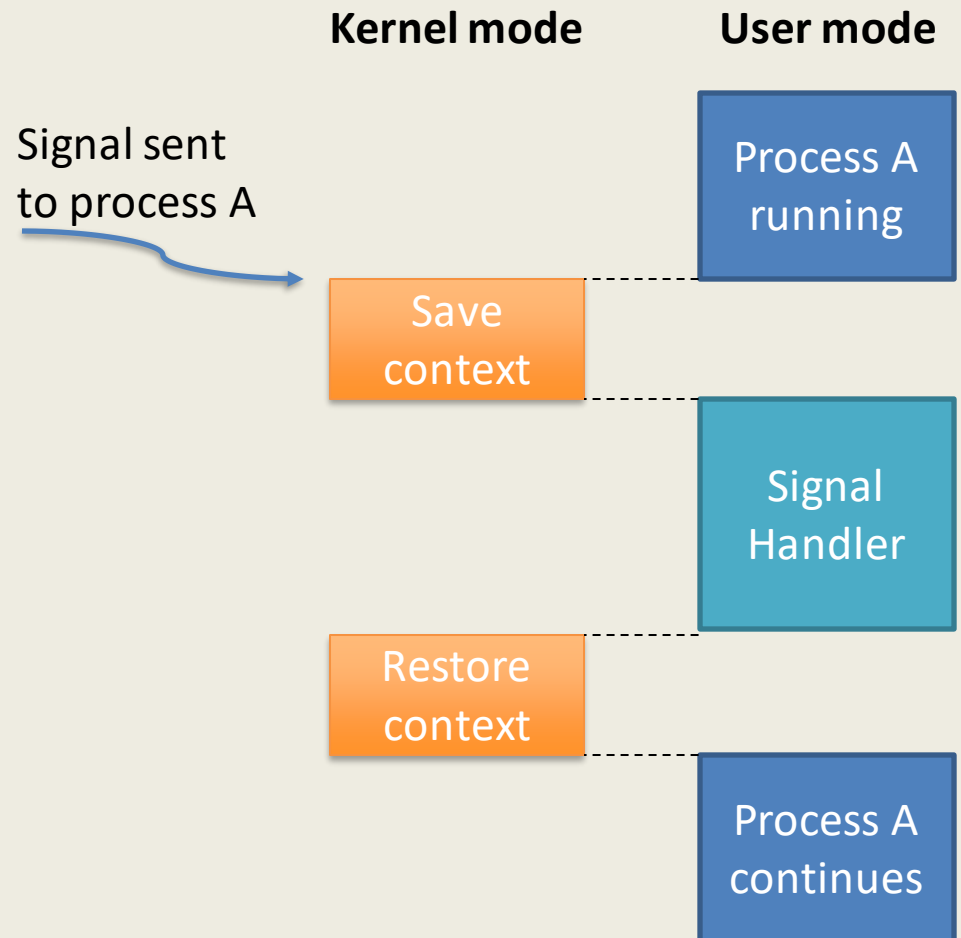
Recalling an interrupt

- ⦿ CPU receives an interrupt
- ⦿ Kernel takes over
 - saves current context
 - handles interrupt
 - restores saved context
- ⦿ Interrupted process continues executing



Handling a signal

- Kernel receives a signal issued to process
- Kernel saves current context
- Process handles signal by calling a specific function
- Kernel restores saved context
- Signaled process continues executing (or maybe terminates)



Which events? Which signals?

- ◉ You can list the signals in your system, with command:

- `$ kill -l`

- ◉ ... and also look at the manual pages:

- `$ man signal`

Some examples of signals

Signal	Default action	Description
SIGINT	Terminate process	Interrupt program
SIGILL	Create core image	Illegal instruction
SIGKILL	Terminate process	Kill program (<i>can't be ignored or handled</i>)
SIGSEGV	Create core image	Segmentation violation (illegal memory access)
SIGALRM	Terminate process	Real-time timer expired
SIGSTOP	Stop process	Stop (<i>can't be ignored or handled</i>)
SIGCONT	Discard signal	Continue after stop
SIGCHLD	Discard signal	Child status has changed
SIGUSR1	Terminate process	User defined signal 1
SIGUSR2	Terminate process	User defined signal 2

Dealing with an incoming signal

- ◉ Default action for specific signal
 - Terminate
 - SIGKILL, SIGINT, SIGUSR1, SIGUSR2
 - Create core image and terminate
 - SIGILL, SIGSEGV, SIGFPE, SIGABORT
 - Stop/Continue
 - SIGSTOP, SIGCONT
 - Nothing (ignore)
 - SIGCHLD, SIGIO, SIGINFO
- ◉ Programmer-defined action (if signal allows)
 - Ignore
 - Handler function

Setting the action for a signal

- ◉ The action for a signal `sig` is set using the `sigaction()` function.

```
#include <signal.h>

int sigaction(int sig,
              const struct sigaction * act,
              struct sigaction * oact);
```

- ◉ `act` sets the action details for signal `sig`, and
- ◉ `oact` (if not `NULL`) is used to store the previously set action details.

Action details

- ◉ `struct sigaction` may differ slightly across architectures but these fields must comply with POSIX.
 - `sa_handler`
 - pointer to an ANSI C handler function
 - `sa_sigaction`
 - pointer to a POSIX handler function
 - `sa_mask`
 - mask of signals to be blocked during signal handling
 - `sa_flags`
 - set of flags that modifies the behaviour of the signal

sa_handler field

- ⦿ Used to set an ANSI C handler function
- ⦿ Possible values:
 - SIG_DFL for the default action
 - SIG_IGN to ignore the signal
 - SIGKILL and SIGSTOP cannot be ignored!
 - `void (*sa_handler) (int)`
(i.e. pointer to an ANSI C signal handler function)

Example: ignoring SIGUSR1

```
/* (...) */

int main(int argc, char *argv[]){
    struct sigaction act;

    /* Clear the act variable. */
    memset(&act, 0, sizeof(struct sigaction));

    act.sa_handler = SIG_IGN;
    sigaction(SIGUSR1, &act, NULL);

    /* SIGUSR1 will now be ignored. */
    /* (...) */
}
```

Example: capturing SIGUSR1

```
/* (...) */

void handle_USR1(int signo)
{
    write(STDOUT_FILENO, "\nCatch USR1!\n", 13);
}

int main(int argc, char *argv[]){
    struct sigaction act;

    memset(&act, 0, sizeof(struct sigaction));
    sigemptyset(&act.sa_mask); /* No signals blocked */
    act.sa_handler = handle_USR1;
    sigaction(SIGUSR1, &act, NULL);

    /* SIGUSR1 will now be captured and handled (ANSI C). */
    /* (...) */
}
```

sa_sigaction field

- ⦿ Used to set a POSIX handler.
 - Provides more power to the signal handler.
 - Requires that the SA_SIGINFO option is set in the sa_flags field.
 - **It cannot be used simultaneously with an ANSI C handler!**
- ⦿ Possible values:
 - ```
void (*sa_sigaction) (int,
 siginfo_t *,
 void *);
```

*(i.e. pointer to a POSIX signal handler function)*

# Example: capturing SIGUSR1

```
void handle_USR1(int signo, siginfo_t *sinfo, void *context)
{
 /* Don't use printf: it is not safe!!! */
 /* See man sigaction for list of safe functions. */
 printf("Signal %d sent by process %d\n",
 sinfo->si_signo, sinfo->si_pid);
}

int main(int argc, char *argv[]) {
 struct sigaction act;

 memset(&act, 0, sizeof(struct sigaction));
 sigemptyset(&act.sa_mask); /* No signals blocked */
 act.sa_sigaction = handle_USR1;
 act.sa_flags = SA_SIGINFO;
 sigaction(SIGUSR1, &act, NULL);

 /* SIGUSR1 will now be captured and handled (POSIX). */
 /* (...) */
}
```



# Signal received during a handler?

- ◉ When a signal is being handled, subsequent notifications of the same signal number are automatically blocked.
  - Those notifications are stored by the OS while the handler is executed.
  - The OS makes them available to be handled afterwards.
- ◉ But what about other signals?
  - By default, they will interrupt the current handler to trigger their own handler.

# sa\_mask field

- ⦿ Used to specify which signals should be blocked during the signal handling.
  - Does not allow an incoming signal to interrupt the signal handling.
- ⦿ A blocked signal is not ignored!
  - The signal is delayed until it becomes unblocked (*i.e. when the handler completes*).
- ⦿ Set is defined on a `sigset_t` variable.
  - Manipulated by sigset operations, check out:
    - `$ man sigsetops`

# Example: setting sa\_mask

```
int main(int argc, char *argv[]){
 struct sigaction act;

 memset(&act, 0, sizeof(struct sigaction));
 sigemptyset(&act.sa_mask); /* No signals blocked */
 sigaddset(&act.sa_mask, SIGINT); /* Block SIGINT */
 sigaddset(&act.sa_mask, SIGUSR2); /* Block SIGUSR2 */

 act.sa_sigaction = handle_USR1;
 act.sa_flags = SA_SIGINFO;
 sigaction(SIGUSR1, &act, NULL);

 /* Now, when a SIGUSR1 is received, SIGINT and SIGUSR2 */
 /* are blocked. */

 /* (...) */

}
```

# Example: setting sa\_mask

```
int main(int argc, char *argv[]){
 struct sigaction act;

 memset(&act, 0, sizeof(struct sigaction));
 sigfillset(&act.sa_mask); /* All signals blocked */
 sigdelset(&act.sa_mask, SIGCHLD); /* Unblock SIGCHLD */
 sigdelset(&act.sa_mask, SIGUSR1); /* Unblock SIGUSR1 */

 act.sa_sigaction = handle_USR2;
 act.sa_flags = SA_SIGINFO;
 sigaction(SIGUSR2, &act, NULL);

 /* Now, when a SIGUSR2 is received, all signals */
 /* except SIGCHLD and SIGUSR1 are blocked. */

 /* (...) */
}
```

# sa\_flags field

- ⦿ Used to modify the behaviour of the signal.
- ⦿ Check out the manual page.
  - `$ man sigaction`

# Signaling from the command line

- Using the kill command, specifying the signal and the target process.

- `$ kill 1234`

- Send (the default) SIGTERM to process with PID 1234.

- `$ kill -SIGKILL 1234`

- Send the SIGKILL to process with PID 1234.

- Check out the manual page:

- `$ man 1 kill`

# Sending a signal from within a process

- ⦿ Using the `kill()` function.
- ⦿ Check out the man page for more information:
  - `$ man 2 kill` (MacOS)
  - `$ man 3 kill` (Linux)

```
#include <signal.h>
```

```
int kill(pid_t pid, int sig);
```

# Example: sending a signal to child

```
int main()
{
 pid_t pid;

 pid = fork();
 if(pid > 0) {
 printf("Hello, my son!\n");
 sleep(5); /* Sleep for 5 seconds... */
 kill(pid, SIGUSR1);
 printf("Goodbye, son!\n");
 }
 else {
 printf("Son is running\n");
 pause(); /* Wait for some signal... */
 }
}
```



# More ways to send signals...

- ⦿ `int raise(int sig) ;`
  - Send the `sig` signal to the current process.
- ⦿ `unsigned int alarm(unsigned int seconds) ;`
  - Send the `SIGALRM` signal to the current process in `seconds` seconds.
- ⦿ `int pause(void) ;`
  - Pauses the current process until a signal is received.
- ⦿ Explore the manual pages.

# Example: setting an alarm

```
void handle_ALARM(int signo)
{
 write(STDOUT_FILENO, "\nTime for a coffee!\n", 20);
}

int main()
{
 struct sigaction act;

 memset(&act, 0, sizeof(struct sigaction));
 sigemptyset(&act.sa_mask); /* No signals blocked */
 act.sa_handler = handle_ALARM;
 sigaction(SIGALRM, &act, NULL);
 /* SIGALRM is now set to be caught and handled. */

 alarm(10);
 /* You have 10 seconds before interrupting the */
 /* following code for a coffee. */
 do_some_very_long_task();
}
```

# Delaying signals

- ⦿ A section of code may require to be immune to a set of signals.
- ⦿ Ignoring the signals during the execution of that section of code was the traditional approach...
  - ... but the notification of the respective event was discarded.
- ⦿ POSIX signals allow to block a set of signals.
  - Signals are stored and received when unblocked.
  - Thus, signals are delayed.
  - Check out the manual page for `sigpending` and `sigprocmask`.

# Example: blocking signals

```
int main()
{
 sigset_t mask, pending;

 sigemptyset(&mask);
 sigaddset(&mask, SIGINT);
 sigprocmask(SIG_BLOCK, &mask, 0);
 /* CTRL-C is now blocked! */

 sleep(10);
 /* Check if SIGINT is pending... */
 sigpending(&pending);
 if(sigismember(&pending, SIGINT))
 printf("SIGINT pending\n");

 sigprocmask(SIG_UNBLOCK, &mask, 0);
 /* CTRL-C is now unblocked! */
}
```