

Shared Memory (1/2)

Sistemas de Computadores
2018/2019


Luís Miguel Pinho / Luis Ferreira / Luís Nogueira / Nuno Pereira

Introduction

- ⦿ Processes are **isolated** one from the other
 - In general it is not possible for one process to access the memory area of another process
- ⦿ There are cases where this restriction is not convenient
- ⦿ It is therefore possible to define a memory area which can be made part of the memory space of two or more processes

Shared memory

- Memory region which is shared by several processes
- Processes with or without an hierarchical link can use this region to communicate



Advantages

- ⦿ Random access
 - Possible to directly access and change any part of the region
- ⦿ Efficiency
 - Accesses are direct (contrarily to pipes, for instance), without the intervention of the OS
 - It is the fastest communication mechanism between processes

Disadvantages

- ◉ Need to synchronize
 - Concurrent processes need to synchronize access to shared memory
- ◉ Pointer sharing
 - We need to assume that the pointers used by one process are **valid only in that particular process**
 - This makes it difficult to implement some data structures (e.g. linked lists or trees)

Introduction

- ◉ In Linux/UNIX there are two APIs of shared memory (SHM)
- ◉ System V (**shmget(2)**)
 - Original mechanism, still very used
 - Part of this course until a few years ago
- ◉ POSIX (**shm_open(3)**)
 - Developed to be more simple to use than older APIs (such as System V)

POSIX Shared Memory in Linux

- ⦿ Implemented as temporary files
 - tmpfs file systems
 - Persistent (exists until explicitly deleted, or system reboot)
- ⦿ A process may map a shared memory area (in POSIX referred frequently as Object), make changes, and disconnect
 - These changes are visible for any process that maps this shared memory object

POSIX Shared Memory in Linux

- ◉ Create

- `shm_open()` , `ftruncate()` , `mmap()`

- ◉ Use the memory

- With a pointer similarly to any dynamic memory

- ◉ Remove (next class)

- `munmap()` , `close()` , `shm_unlink()`

POSIX Shared Memory in Linux

- ◉ Usual headers

- `<fcntl.h>` - file control options
- `<sys/types.h>` - data types used by the API
- `<sys/stat.h>` - constants used for opening
- `<sys/mman.h>` - declarations related with management of shared memory : **shm_*** and **mmap ()**

POSIX Shared Memory in Linux

- ⦿ Compiling

- In Linux, the functions are in what is called “POSIX realtime library (librt)”, so it is necessary to compile and link with the library (rt)

- ⦿ Exemplo:

```
gcc -Wall -o shm_test shm_test.c -lrt
```

Create

1. **shm_open ()** : Creates and opens a shared memory area
 - If already exists, just opens to use
 - Returns a file descriptor to be used in other functions
2. **ftruncate ()** : Defines the size
3. **mmap ()** : Maps the shared area in the process address space

shm_open ()

```
#include <sys/mman.h>
#include <sys/stat.h>      /* For constants "mode" */
#include <fcntl.h>         /* For constants O_* */

int shm_open(const char *name, int oflag, mode_t mode );
```

- ⦿ Creates and opens a shared memory area, or opens an existing one
- ⦿ Returns a file descriptor (the shared memory), or -1 in case of error

shm_open()

```
#include <sys/mman.h>
#include <sys/stat.h>      /* For constants "mode" */
#include <fcntl.h>         /* For constants O_* */

int shm_open(const char *name, int oflag, mode_t mode );
```

- ⦿ **name**: name to be used to identify the area
 - Must start with '/'

shm_open()

```
#include <sys/mman.h>
#include <sys/stat.h>      /* For constants "mode" */
#include <fcntl.h>          /* For constants O_* */

int shm_open(const char *name, int oflag, mode_t mode );
```

⦿ **oflag**: Creation options

- **O_CREAT**: creates, if not existing, or use existing
 - Without **O_CREAT** tries to use existing (error if none)
- **O_EXCL**: (with **O_CREAT**) **just creates**, error if already exists
- **O_TRUNC**: eliminates contents already existing in the shared memory area

⦿ Access Flags

- **O_RDONLY**: open for reading only
- **O_RDWR**: open for reading and writing (no flag to just write)

shm_open ()

```
#include <sys/mman.h>
#include <sys/stat.h>      /* For constants "mode" */
#include <fcntl.h>         /* For constants O_* */

int shm_open(const char *name, int oflag, mode_t mode );
```

- ⦿ **mode**: Defines permissions, as per Linux usual style
 - RWX for user/group/others
 - Makes AND with the complement of the process umask
 - Should be 0 (zero) if opening an existing area

ftruncate()

```
#include <unistd.h>
#include <sys/types.h>

int ftruncate(int fd, off_t length);
```

- ⦿ Defines the size of the area and initializes it to zero
 - Must be used after **shm_open()**, as memory is created with size zero
 - If area already exists, only changes if size is different
 - If new size > than existing, the size is increased with zero in the new area
 - If new size < than existing, size is reduced (losing the extra data)
- ⦿ Returns 0 if successful, or -1 in case of error (*errno* has the error)

ftruncate()

```
#include <unistd.h>
#include <sys/types.h>

int ftruncate(int fd, off_t length);
```

- ⦿ **fd:**

- File descriptor (returned by **shm_open()**)

- ⦿ **length:**

- size, in bytes, of memory object

mmap ()

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int  
flags, int fd, off_t offset);
```

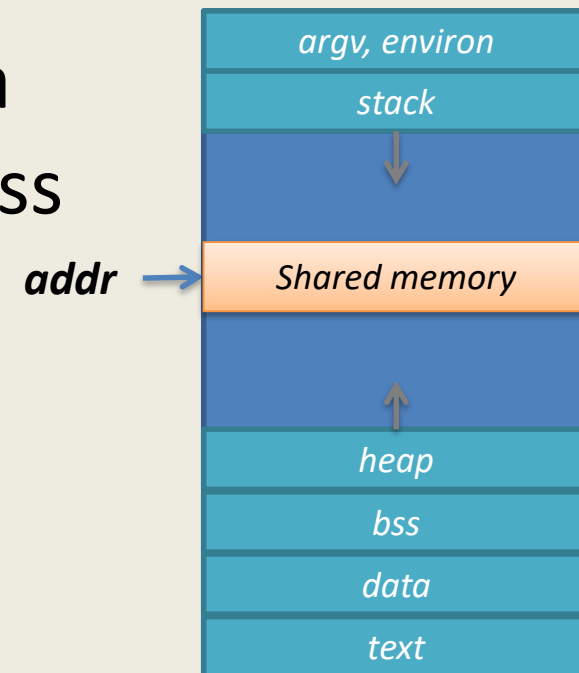
- ⦿ Maps a shared memory object in the process address space
 - In practice, we are required to handle arguments ***length*** and ***prot***
- ⦿ Returns a pointer to the object or **MAP_FAILED ((void *) -1)**, with *errno* set with the error

mmap ()

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int  
          flags, int fd, off_t offset);
```

- ⦿ **addr**: allows to request the area to be mapped in a specific address
 - **NULL** in the usual case to let the OS define the address



mmap ()

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int  
flags, int fd, off_t offset);
```

- ⦿ **length**: Required size to be used
 - \leq than size defined by **ftruncate ()**
 - OS usually rounds to a multiple of the size of a memory page

mmap ()

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int  
flags, int fd, off_t offset);
```

- ◉ **prot**: protection flags

- **PROT_READ**: just read
- **PROT_READ | PROT_WRITE**: read and write

- ◉ Must be consistent with **shm_open ()**

- Do not use **O_RDONLY** in **shm_open ()** and then **PROT_READ | PROT_WRITE** in **mmap ()**

mmap ()

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int  
    flags, int fd, off_t offset);
```

⦿ **flags**: control behaviour of **mmap ()**

- In the case of shared memory, always use **MAP_SHARED**, to allow changes be seen by other processes

⦿ **fd**: File descriptor

- Returned by **shm_open ()**

mmap ()

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int  
flags, int fd, off_t offset);
```

⦿ **offset**: beginning of mapping in the memory area

- Usually zero, to use all memory area from the beginning
- Must be a multiple of memory page

Create and open

```
1.int fd;  
2.void * addr;  
3.fd = shm_open("/shmtest", O_CREAT|O_EXCL|  
   O_RDWR, S_IRUSR|S_IWUSR);  
4.ftruncate (fd, 100);  
5.addr = mmap(NULL, 100, PROT_READ|PROT_WRITE,  
   MAP_SHARED, fd, 0);
```

Line 3: Creates (**O_CREAT**) shared memory area (with error if existis because of **O_EXCL**) with name **"/shmtest"**, and read/write permissons (**O_RDWR**), and open for user to read (**S_IRUSR**) and write (**S_IWUSR**).

Criar e Abrir Memória Partilhada

```
1.int fd;  
2.void * addr;  
3.fd = shm_open("/shmtest", O_CREAT|O_EXCL|  
    O_RDWR, S_IRUSR|S_IWUSR);  
4.ftruncate (fd, 100);  
5.addr = mmap(NULL, 100, PROT_READ|PROT_WRITE,  
    MAP_SHARED, fd, 0);
```

Line 4: Defines a 100 bytes size

Criar e Abrir Memória Partilhada

```
1.int fd;  
2.void * addr;  
3.fd = shm_open("/shmtest", O_CREAT|O_EXCL|  
    O_RDWR, S_IRUSR|S_IWUSR);  
4.ftruncate (fd, 100);  
5.addr = mmap(NULL, 100, PROT_READ|PROT_WRITE,  
    MAP_SHARED, fd, 0);
```

Line 5: Maps the area in na address decided by the OS (**NULL**), requiring to use all the size (**100**), with read and write permissions (**PROT_READ | PROT_WRITE**), consistente with open (continues)

Criar e Abrir Memória Partilhada

```
1.int fd;  
2.void * addr;  
3.fd = shm_open("/shmtest", O_CREAT|O_EXCL|  
    O_RDWR, S_IRUSR|S_IWUSR);  
4.ftruncate (fd, 100);  
5.addr = mmap(NULL, 100, PROT_READ|PROT_WRITE,  
    MAP_SHARED, fd, 0);
```

Line 5: (continued) Area is shared (**MAP_SHARED**), and it refers to the previous open (**fd**). All memory área will be used so starting from beginning (offset is 0).

Use

- ⦿ Use as normal dynamic memory
 - The pointer obtained with **mmap ()** can be used the same way as pointers given by **malloc ()**
 - It is common to have a structure defined for the data in the memory area
- ⦿ Next example
 - Shared memory between unrelated processes (no parent - child)
 - Reading / writing in a data structure
 - No synchronization - writer needs to be executed first

Escritor

```
1. typedef struct {
2.     int var1;
3.     int var2;
4. } shared_data_type;

5. int main(int argc, char *argv[]) {
6.     int fd, data_size = sizeof(shared_data_type);
7.     shared_data_type *shared_data;
8.     int fd = shm_open("/shmtest", O_CREAT|O_EXCL|O_RDWR,
9.         S_IRUSR|S_IWUSR);
10.    ftruncate (fd, data_size);
11.    shared_data = (shared_data_type*)mmap(NULL,
12.        data_size, PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0);
13.    shared_data->var1 = 100;
14.    shared_data->var2 = 200;
15.}
```

Write

Leitor

```
1. typedef struct {
2.     int var1;
3.     int var2;
4. } shared_data_type;

5. int main(int argc, char *argv[]) {
6.     int fd, data_size = sizeof(shared_data_type);
7.     shared_data_type *shared_data;
8.     int fd = shm_open("/shmtest", O_RDWR,
9.         S_IRUSR|S_IWUSR);
10.    ftruncate (fd, data_size);
11.    shared_data = (shared_data_type *)mmap(NULL,
12.        data_size, PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0);
13.    printf("%d\n", shared_data->var1);
14.    printf("%d\n", shared_data->var2);
15. }
```

No O_CREAT

Read

Use

- New example with synchronization with active wait (spinning)
 - Processes are related
 - Parent is **writer** and son is **reader**

```
1. typedef struct {
2.     int var1;
3.     int var2;
4.     int new_data;
5. } shared_data_type;

6. int main(int argc, char *argv[]) {
7.     int fd, data_size = sizeof(shared_data_type);
8.     shared_data_type *shared_data;
9.     fd = shm_open("/shmtest", O_CREAT|O_EXCL|O_RDWR,
10.         S_IRUSR|S_IWUSR);
11.     ftruncate (fd, data_size);
12.     shared_data = (shared_data_type *)mmap(NULL,
13.         data_size, PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0);
```

Utilizar a Memória Partilhada

```
7.  shared_data->new_data = 0;
8.  pid=fork()
9.  if (pid > 0) { /* writer*/
10.     shared_data->var1 = 100;
11.     shared_data->var2 = 200;
12.     shared_data->new_data = 1;
13.     wait();
14. }
15. else { /* reader*/
16.     while(!shared_data->new_data);
17.     printf("%d\n", shared_data->var1);
18.     printf("%d\n", shared_data->var2);
19. }
20. }
```

Write (Parent)

Synchronization

Read (Son)

Exercise TP5.1

- ⦿ Create two unrelated processes:
 - Writer: writes string in shared memory
 - Reader: read and print the message.
- ⦿ Reader executed after writer

Exercise TP5.2

○ Number 5 lotary:

- Writer: Places random numbers ¹ [1,5] in shared memory.
- Reader: waits for a new number (spinning in a flag), and prints the number.
- Must guarantee that writer waits for reader to finish processing a number before generating a new one (writer also spins in a flag)
- Both finish when number is 5. Reader prints how many numbers received.

¹ Random numbers in an interval [1,N] can be generated with:

```
#include <stdio.h>
#include <stdlib.h>
...
/* Generator seed */
srand((unsigned) time(NULL));
...
/* generate number */
num = rand() % N + 1;
```

Solution

```
1. typedef struct {
2.     int loto;
3.     int canRead;
4.     int canWrite;
5. } shared_data_type;

6. int main(int argc, char *argv[]) {
7.     int fd, data_size = sizeof(shared_data_type);
8.     shared_data_type *shared_data;
9.     fd = shm_open("/shmtest", O_CREAT|O_EXCL|O_RDWR,
10.         S_IRUSR|S_IWUSR);
11.     ftruncate (fd, data_size);
12.     shared_data = (shared_data_type *)mmap(NULL,
13.         data_size, PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0
14.     );
```

Writer process

```
7.  shared_data->canRead = 0;
8.  shared_data->canWrite = 1;
9.  pid=fork()
10. if (pid > 0) {
11.     randomNum = 0;
12.     while (randomNum != 5) {
13.         randomNum = generateRandom();
14.         while(shared_data->canWrite == 0);
15.         shared_data->canWrite = 0;
16.         shared_data->loto = randomNum;
17.         shared_data->canRead = 1;
18.     }
19. Else {
```

Reader process

```
7.  randomNum = 0
8.  while (randomNum != 5)
9.  {
10.     while(shared_data->canRead == 0);
11.     shared_data->canRead = 0;
12.     randomNum = shared_data->loto;
13.     shared_data->canWrite = 1;
14. }
```