

Shared Memory (2/2)


Sistemas de Computadores
2017/2018

Luís Miguel Pinho / Luis Ferreira / Luís Nogueira / Nuno Pereira

Previous class: POSIX shared mem in Linux

- Shared memory

- Allow two or more processes to access the same memory area



- Create

- `shm_open()` , `ftruncate()` , `mmap()`

- Use

- With a pointer as dynamic memory

- This class: remove

- `munmap()` , `close()` , `shm_unlink()`

Remove

- ◉ **`munmap ()`**

- Disconnects the shared memory area from the process address space (inverse of **`mmap ()`**)

- ◉ **`close ()`**

- Closes the file descriptor returned by **`shm_open ()`**

- ◉ **`shm_unlink ()`**

- Removes the memory area from the file system. Marks it to be deleted, as soon as all processes using it close it

`munmap ()`

```
#include <sys/mman.h>
```

```
int munmap(void *addr, size_t length);
```

- ⦿ Disconnects the shared memory area from the process address space
- ⦿ Returns 0 if successful, or -1 in case of error (*errno* is set with the error)

munmap ()

```
#include <sys/mman.h>
```

```
int munmap(void *addr, size_t length);
```

- ⦿ **addr**: Pointer returned by **mmap()**
- ⦿ **length**: Size of the mapped memory area
 - \leq size defined in **ftruncate()**
 - Usually rounded to memory page multiple

close ()

```
#include <unistd.h>
```

```
int close(int fd);
```

- ⦿ Closes the file descriptor
- ⦿ **fd**: the file descriptor returned by **shm_open ()**
- ⦿ Returns 0 if successful, or -1 in case of error (*errno* is set with the error)

shm_unlink()

```
#include <sys/mman.h>
```

```
int shm_unlink(const char *name);
```

- ◉ Removes memory area from file system
 - May not be immediate if any process still has it open
 - The area is marked to be removed as soon as all processes close it
- ◉ **name**: shared memory name
- ◉ Returns 0 if successful, or -1 in case of error (*errno* is set with the error)

Close shared memory

- ◉ Example:

- Process close its descriptor as soon as not needing it
- Another process can still open and use it

Close shared memory

```
1.  int fd, r;
2.  void *addr; /* pointer to the shared memory */
3.  fd = shm_open("/shmtest", ...);
4.  ftruncate (fd, 100);
5.  addr = mmap(NULL, 100, ...);

6.  ... /* use - read/write */

7.  r = munmap(addr, 100); /* disconnects */
8.  if (r < 0) exit(1); /* Check error */

9.  r = close(fd); /* closes */
10. if (r < 0) exit(1); /* Check error*/
11. exit(0); /* all descriptors are closed... */
```

Remove shared memory

- ◉ Next example:
 - Process removes the shared memory
 - Other processes will have to create it again

Remove shared memory

```
1.  int fd, r;
2.  void *addr; /* pointer to the shared memory */
3.  fd = shm_open("/shmtest", ...);
4.  ftruncate (fd, 100);
5.  addr = mmap(NULL, 100, ...);

6.  ... /* use - read/write */

7.  r = munmap(addr, 100); /* disconnects */
8.  if (r < 0) exit(1); /* Check error */

9.  r = shm_unlink("/shmtest"); /* removes */
10. if (r < 0) exit(1); /* Check error */
11. exit(0);
```

Exercise TP6.1

- ⦿ Create two unrelated processes: writer and reader.
 - Writer: asks user the name of the place and a set of 10 temperature readings of the last 24 hours, putting all this information in the shared memory
 - Leitor: Calculates and prints average
- ⦿ The two processes may be executed simultaneously
- ⦿ Reader removes area from the system

Exercise TP6.2

- ◉ Implement a program that creates a shared memory area for two arrays of integer: **v[1000]** e **max[10]**
 - Initialize **v[]** with random numbers
- ◉ Create 10 children; Each child should search the max in 1/10 of **v[]** and put it in **max[i]** (**i** is the child index)
- ◉ Parents waits all children to terminate and calculates the maximum in **max[]**
- ◉ All processes should disconnect and close
- ◉ Parent should remove shared memory area before terminating

Obtain information about shared memory

```
#include <sys/stat.h>

int fstat(int fd, struct stat *buf);
```

- ◉ Args:
 - **fd**: file descriptor returned **shm_open()**
 - **buf**: **stat** structure where information will be placed
- ◉ The structure will have info on the area, such as size, permissions, owner, etc.

Obtain information about shared memory

```
struct stat {  
    dev_t      st_dev;      /* ID of device containing file */  
    ino_t      st_ino;      /* inode number */  
    mode_t     st_mode;     /* protection */  
    nlink_t    st_nlink;    /* number of hard links */  
    uid_t      st_uid;      /* user ID of owner */  
    gid_t      st_gid;      /* group ID of owner */  
    dev_t      st_rdev;     /* device ID (if special file) */  
    off_t      st_size;     /* total size, in bytes */  
    blksize_t  st_blksize;  /* blocksize for file system I/O */  
    blkcnt_t   st_blocks;   /* number of 512B blocks allocated */  
    time_t     st_atime;    /* time of last access */  
    time_t     st_mtime;    /* time of last modification */  
    time_t     st_ctime;    /* time of last status change */  
};
```

Obtain information about shared memory

```
int main(void){
    int fd; /* File descriptor */
    int ret;
    struct stat shm_stat;

    fd = shm_open("/shmtest", O_CREAT|O_EXCL|O_RDWR, 0600);
    if (fd < 0) exit(1);

    ftruncate (fd, 100);
    addr = mmap(NULL,100,PROT_READ|PROT_WRITE,MAP_SHARED,fd,0);
    if (addr == NULL) exit(1);

    if(fstat(fd, &shm_stat) < 0) exit (1);

    printf("mode = %d\n", shm_stat.st_mode);
    printf("size = %d\n", shm_stat.st_size);
    ...
}
```


/dev/shm

- ◉ In Linux, the temporary file system (`tmpfs`) used for POSIX shared memory is mapped in **/dev/shm**
 - POSIX semaphores (next classes) are also mapped here
- ◉ Shared memory objects can be listed as files (with **ls**)

```
% ls /dev/shm  
shmtest  shmtp61
```

/dev/shm

- ◉ And removed with **rm**

```
% rm /dev/shm/shmtest
```